



# Private Cloud Test-Drive

---

## Test-Drive Guide

(Cloud Infrastructure)

**Prepared by:**

**Roberto Calva - LATAM Management & Automation Specialist SA**

---

**Vagner Farias - LATAM Cloud Infrastructure Specialist SA**

**Creation date:** 06/28/18

**Date of last modification:** 06/02/19

**Version:** 1.1

---

# Lab 1: Self Service VM Creation


This first lab aims to provide an overview of the OpenStack Dashboard (Horizon) and you'll learn to:

- Create virtual machines (commonly called virtual instances in cloud environments) connected to an existing network.
- Access the console of a virtual instance.

Details on how to manage projects, users, networks and volumes will be covered in subsequent labs.

## OpenStack Dashboard Overview

To access the OpenStack Dashboard, use the URL provided to you (eg. <https://openstack-GUID.rhpds.opentlc.com>).



Use the following credentials to login:


Login	td-user
Password	r3dh4t1!

This user will allow you to access a project called “test-drive” with a “\_member\_” role. This essentially means that this user is not an OpenStack administrator and is unable to

manage user permissions, create new projects or manage quotas, among other administrative tasks.

Inside the “test-drive” project, however, this user can manage virtual instances, networks, routers, images, volumes.

First thing you may do is to explore the dashboard. In the top left section of the screen you have two tabs: Project and Identity.



As you’re logged as a “\_member\_”, the *Identity* tab will only show information about your user. OpenStack administrators can manage users and projects.

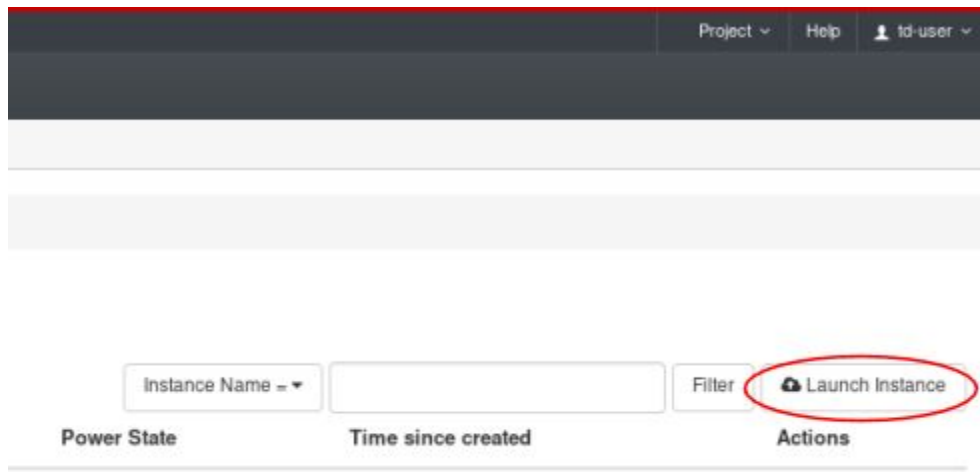
In the Project tab you can access several other sections that will allow you to execute several tasks in the environment, as summarized in the following table.

Section	Description
Compute	Provides an overview of used and available resources, considering user quota. It’s also the section where you’ll be able to create new instances, volumes, images or key pairs as well as manage existing ones.
Network	Provides a network topology diagram and allows you to manage networks, routers, security groups and floating IPs.
Orchestration	Provides an interface to the Orchestration service (Heat), which as the name implies allows orchestrating the creation/association of OpenStack resources to deploy applications based on a template.

At the top right corner of the screen you would be able to switch between projects if there were other projects you had access to, which is not the case at this moment. Also, clicking in your username you have the option to change some User Interface settings or your password. It’s also through this menu that you may logout from the dashboard.

## Create a new instance

To create a new instance go to Project → Compute → Instances and click on the “Launch Instance” button as shown in the diagram below.



The Launch Instance window will be opened and you need to provide at least the information marked with a “\*”. For this lab, you should use the following information:

Instance name	instance1
Availability Zone	nova
Count	1
Select Boot Source	Image
Create New Volume	No
Image	webserver
Flavor	m1.small
Network	private

The following images demonstrate the process.

### Details tab

## Launch Instance



### Details

Source \*

Flavor \*

Networks

Network Ports

Security Groups

Key Pair

Configuration

Server Groups

Scheduler Hints

Metadata

Please provide the initial hostname for the instance, the availability zone where it will be deployed, and the instance count. Increase the Count to create multiple instances with the same settings.



Instance Name \*

instance1

Availability Zone

nova

Count \*

1

Total Instances (10 Max)



0 Current Usage  
1 Added  
9 Remaining

✕ Cancel

< Back

Next >

Launch Instance

## Source tab

## Launch Instance



Details

Instance source is the template used to create an instance. You can use a snapshot of an existing instance, an image, or a volume (if enabled). You can also choose to use persistent storage by creating a new volume.



Source

Select Boot Source

Image

Create New Volume

Yes No

Flavor \*

Allocated

Name	Updated	Size	Type	Visibility	
> webserver	8/24/18 4:57 PM	601.13 MB	qcow2	Public	-

Available 2

Select one

Click here for filters.

Name	Updated	Size	Type	Visibility	
> RHEL7	7/25/18 9:45 PM	546.84 MB	qcow2	Public	+
> cirros	7/17/18 1:03 AM	12.65 MB	qcow2	Public	+

Cancel

< Back

Next >

Launch Instance

## Flavor tab

## Launch Instance



Details

Flavors manage the sizing for the compute, memory and storage capacity of the instance.



Source

Allocated

Name	VCPUS	RAM	Total Disk	Root Disk	Ephemeral Disk	Public	
> m1.small	1	2 GB	20 GB	20 GB	0 GB	Yes	-

Available 1

Select one

Click here for filters.

Name	VCPUS	RAM	Total Disk	Root Disk	Ephemeral Disk	Public	
> m1.tiny	1	512 MB	1 GB	1 GB	0 GB	Yes	+

Cancel

< Back

Next >

Launch Instance

## Networks tab

Launch Instance

Details

Source

Flavor

**Networks**

Network Ports

Security Groups

Key Pair

Configuration

Server Groups

Scheduler Hints

Metadata

Networks provide the communication channels for instances in the cloud.

▼ Allocated 1

Select networks from those listed below.

Network	Subnets Associated	Shared	Admin State	Status
> private	private-subnet	No	Up	Active

▼ Available 1

Select at least one network

Click here for filters.

Network	Subnets Associated	Shared	Admin State	Status
> external	external-subnet	Yes	Up	Active

Cancel

< Back

Next >

Launch Instance

After the instance creation finishes you'll be able to see it in the listing at Project → Compute → Instances. At this point, you can use the Actions menu and access the instance console.



To log into this instance, use:

Login	root
Password	r3dh4t1!

Test connectivity to the virtual router (10.10.10.1) and to an Internet address (eg. 8.8.8.8)



```

[root@instance1 ~]# ping -c3 10.10.10.1
PING 10.10.10.1 (10.10.10.1) 56(84) bytes of data.
64 bytes from 10.10.10.1: icmp_seq=1 ttl=64 time=1.64 ms
64 bytes from 10.10.10.1: icmp_seq=2 ttl=64 time=0.768 ms
64 bytes from 10.10.10.1: icmp_seq=3 ttl=64 time=0.489 ms

--- 10.10.10.1 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2002ms
rtt min/avg/max/mdev = 0.489/0.966/1.641/0.490 ms
[root@instance1 ~]#
[root@instance1 ~]# ping -c3 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=121 time=4.84 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=121 time=1.82 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=121 time=1.96 ms

--- 8.8.8.8 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2004ms
rtt min/avg/max/mdev = 1.821/2.876/4.844/1.393 ms
[root@instance1 ~]#

```

You can check the network topology at Project → Network → Network Topology.

Project / Network / Network Topology


## Network Topology

Topology

Graph

Small

Normal



## Create a second instance

In order to test connectivity between two instances, launch a new one (Project → Compute → Instances → Launch Instance), using the following information:

Instance name	Instance2
Availability Zone	nova
Count	1
Select Boot Source	Image
Create New Volume	No
Image	webserver
Flavor	m1.small
Network	private

From the Instances list, take note of the IP address of **instance1**. After that, connect to **instance2** console, log into it and ping the IP address of **instance1** and the router (10.10.10.1).

## Lab 2: External Inbound Access to Instances

In this lab you'll learn:

- How to use floating IPs to allow inbound access to virtual instances.

### Virtual Instance IP addresses

In order to allow inbound access to a virtual instance connected to a virtual private network it's required to associate a floating IP to it. At this time, it's important to differentiate Fixed IPs and Floating IPs.

#### Fixed IP Address

When an instance is created it receives at least one IP address, which is known as the *fixed IP* and is provided by DHCP<sup>1</sup> according to the network to which the virtual NIC is connected to. The number of fixed IPs the instance will have depends on the number of virtual NICs configured.

This IP address is visible from within the instance (eg. you can see it running “ip address show”) and is typically part of a virtual private network.

#### Floating IP Address

Floating IP addresses are used to allow external inbound access to instances connected to virtual private networks. The floating IPs are created on an *external network* and when it's attached to a virtual instance, neutron will configure a Destination NAT (DNAT) from the external IP address to the fixed IP address used by the instance.

Virtual instances don't have any information about floating IP addresses and the address translation will be handled by neutron L3 agent (in the case of lab environment).

### Associate Floating IP to instance

1. Log in as td-user.
2. Go to Projects → Compute → Instances.

---

<sup>1</sup> Actually IPv6 may use SLAAC to configure IP addresses.

3. In “instance1” Actions menu, select “Associate Floating IP”



4. In the next dialog, select one of the available IP addresses and click on “Associate”

A screenshot of the 'Manage Floating IP Associations' dialog box. The dialog has a title bar with a close button (X). It contains two main sections. The first section is labeled 'IP Address' with a blue asterisk. It features a dropdown menu showing '192.168.200.108' and a plus sign button. To the right of this section is a text label: 'Select the IP address you wish to associate with the selected instance or port.' The second section is labeled 'Port to be associated' with a blue asterisk. It features a dropdown menu showing 'instance1: 10.10.10.106'. At the bottom right of the dialog are two buttons: 'Cancel' and 'Associate'.

“Port to be associated” sets to which IP address the DNAT configuration should be done. In this case, the instance has only one port/IP address.

5. Repeat the process for *instance2*.



Only Floating IP addresses **192.168.200.106** and **192.168.200.108** will be externally accessible.

When you finish this procedure, you'll be able to see the floating IP associated to the instance, in addition to its fixed IP (Project → Compute → Instances):

<input type="checkbox"/>	Instance Name	Image Name	IP Address
<input type="checkbox"/>	instance1	webserver	<div><ul style="list-style-type: none"><li>• 10.10.10.106</li></ul>Floating IPs:<ul style="list-style-type: none"><li>• 192.168.200.108</li></ul></div>



Take note of the floating IP associated to your instance as you'll use it in the next step.

## Check Instance IP Addresses

To confirm the floating IP address is unknown by the virtual instance, access the instance console using the following steps:

1. Through the OpenStack Dashboard, go to Projects → Compute → Instances.
2. From the instance “Actions” menu, select “Console”.
3. Run “ip address show” and check only the fixed IP address is known by the instance.

```
[root@instance1 ~]# ip address show
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN qlen 1
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1446 qdisc pfifo_fast state UP qlen 1000
    link/ether fa:16:3e:9b:00:d6 brd ff:ff:ff:ff:ff:ff
    inet 10.10.10.106/24 brd 10.10.10.255 scope global dynamic eth0
        valid_lft 77327sec preferred_lft 77327sec
    inet6 fe80::f816:3eff:fe9b:d6/64 scope link
        valid_lft forever preferred_lft forever
[root@instance1 ~]# _
```

## Test External Inbound Access

In this environment, each of the two pre-configured IP addresses is associated to a resolvable hostname. Using other floating IP addresses won't work as these are the only ones with mapping configured.

192.168.200.106	floating106-GUID.rhpds.opentlc.com
192.168.200.108	floating108-GUID.rhpds.opentlc.com

Although the instance was configured with a floating IP, it won't be possible to access any service on it as the OpenStack security groups are preventing such access. This will be covered on the next lab.

## Lab 3: Security Groups

In this lab you'll learn:

- What are security groups.
- How to use security groups to allow or block access to instances.

### Overview of Security Groups

In a nutshell, security groups act as a firewall at the instance level, controlling both inbound and outbound access.

In this case, “at the instance level” doesn’t mean to be running a firewall *inside* the virtual instance operating system which could cause too much overhead at scale, but to apply rules at the virtual switch level. This means there’s no need to configure firewall rules inside the virtual instance.

In the reference implementation of security groups, used in this environment and in production clouds, the rules are applied using *iptables* in the compute nodes, directly on the virtual switch ports used by the virtual instance.

This is different from the most common approach in legacy infrastructures where there’s a firewall in the perimeter in addition to network switch Access Control Lists (ACLs) controlling access to and from devices.

Information needed to configure security groups are the following:

- Source / destination
- Protocol
- Port

### Check Security Groups for Virtual Instance

It’s time to understand why you could successfully get *ping* replies from *instance1* but couldn’t connect to it using *ssh*.

1. Log into the OpenStack Dashboard as *td-user*.
2. Go to Project → Compute → Instances, click on *instance1*.

3. At the *Overview* tab, look at the Security Groups section, as illustrated below

Security Groups	
default	ALLOW IPv4 to 0.0.0.0/0 ALLOW IPv4 from default ALLOW IPv4 icmp from 0.0.0.0/0 ALLOW IPv6 to ::/0 ALLOW IPv6 from default

The highlighted rule allowed the instance to reply to *ping* from any host. Examining each rule in the security group:

ALLOW IPv4 to 0.0.0.0/0	Allows <i>egress</i> traffic from the instance <i>to</i> any host, to any port, using IPv4.
ALLOW IPv4 from default	Allows <i>ingress</i> traffic from any instance that also has the <i>default</i> security group applied, using IPv4.
ALLOW IPv4 icmp from 0.0.0.0/0	Allows <i>ingress</i> traffic of <i>icmp</i> protocol <i>from</i> any host, using IPv4.
ALLOW IPv6 to ::/0	Allows <i>egress</i> traffic from the instance <i>to</i> any host, to any port, using IPv6.
ALLOW IPv6 from default	Allows <i>ingress</i> traffic from any instance that also has the <i>default</i> security group applied, using IPv6.

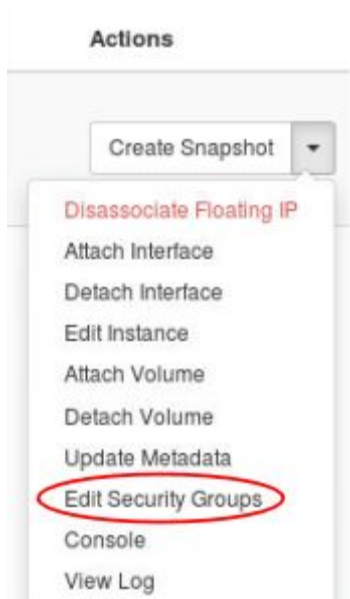
## Apply Security Group to Virtual Instance

To allow HTTP and SSH access to the instances, apply the security group SSH\_HTTP\_From\_any to them, executing the following steps:

1. Log into the OpenStack Dashboard as *td-user*.
2. Go to Project → Compute → Instances



3. In the “Actions” menu for the instance, click on “Edit Security Groups”.



4. In the following dialog, you’ll be able to see all available security groups on the left and all applied security groups on the right. Click on the “+” sign for the “SSH\_HTTP\_from\_any” security group to add it to the instance.



5. Check applied security groups going to the Overview tab for the instance (Project → Compute → Instances → Instance1 → Overview).


SSH_HTTP_From_Any	<ul style="list-style-type: none"><li>• ALLOW IPv6 to ::/0</li><li>• ALLOW IPv4 22/tcp from 0.0.0.0/0</li><li>• ALLOW IPv4 8080/tcp from 0.0.0.0/0</li><li>• ALLOW IPv4 443/tcp from 0.0.0.0/0</li><li>• ALLOW IPv4 to 0.0.0.0/0</li><li>• ALLOW IPv4 80/tcp from 0.0.0.0/0</li><li>• ALLOW IPv4 9990/tcp from 0.0.0.0/0</li></ul>
default	<ul style="list-style-type: none"><li>• ALLOW IPv4 icmp from 0.0.0.0/0</li><li>• ALLOW IPv6 to ::/0</li><li>• ALLOW IPv6 from default</li><li>• ALLOW IPv4 to 0.0.0.0/0</li><li>• ALLOW IPv4 from default</li></ul>

## Test HTTP Access From External Host

Identify the floating IP address and access the the corresponding URL in your web browser.

192.168.200.106	<a href="http://floating106-GUID.rhpdn.opentlc.com">http://floating106-GUID.rhpdn.opentlc.com</a>
192.168.200.108	<a href="http://floating108-GUID.rhpdn.opentlc.com">http://floating108-GUID.rhpdn.opentlc.com</a>

Example:



This is web server at 10.10.10.106

## Lab 4: Persistent Volume Management

In this lab you'll learn:

- What are volumes.
- How to use volumes to manage data life cycle.

### Overview of Volumes

A volume is a persistent block storage device that can be attached to or detached from instances. Currently, volumes can only be attached to one instance at a time.

Without volumes, instances only have ephemeral storage provided by the flavor configuration, meaning that all data is lost when an instance are deleted.

Besides, ephemeral storage commonly use local compute nodes disks, so if the node is unavailable, data is unavailable as well. An exception to this is when OpenStack Compute service (nova) is configured to use a shared storage solution, like NFS or Red Hat Ceph Storage.

Volumes can be used as the primary disk for the instance (where the operating system resides) or secondary disks used for storing data.

Volumes are provided by OpenStack Block Storage service, widely known by its project name *cinder*.

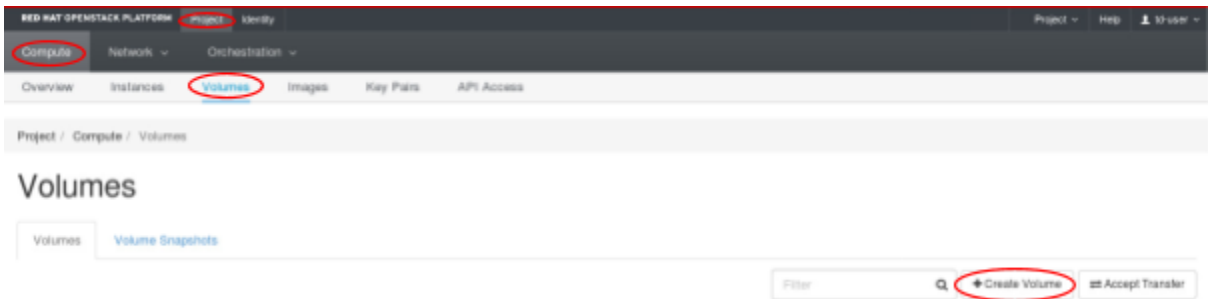
The OpenStack Block Storage service creates an abstraction layer between the real storage device in use and the instances. No matter if the storage device is a SAN, a NAS or a scale-out storage solution like Red Hat Ceph Storage, the operating system inside the instances will see a standard block device.

### Volumes for Secondary Disks

Suppose you have an application that needs to store data and that you need to ensure this data won't be lost or made unavailable if there are issues with the compute node. In this case, you can keep the operating system in ephemeral disk and use a secondary disk for data. To experiment, do the following:

1. Log into the OpenStack Dashboard as td-user.

2. Go to Project → Compute → Volumes and click on “+Create Volume” button.



3. Create a volume with the following configuration.

Volume Name	volume1
Description	Data volume
Volume Source	No source, empty volume
Volume Type	standard
Size (GB)	1
Availability Zone	nova

### Create Volume

Volume Name

volume1

Description

Data Volume

Volume Source

No source, empty volume

Type

standard

Size (GiB) \*

1

Availability Zone

nova

Description:

Volumes are block devices that can be attached to instances.

Volume Type Description:

**standard**  
No description available.

Volume Limits

Total GiBbytes

0 of 1,000 GiB Used


Number of Volumes

0 of 10 Used

Cancel

Create Volume

4. From the volume list (Project → Compute → Volumes), select the “Manage Attachments” option in Actions Menu for *volume1*.



5. In the following dialog, select *instance1* in the “Attach to Instance” drop down menu.

**Manage Volume Attachments** ✕

Instance	Device	Actions
No items to display.		

**Attach To Instance**

Attach to Instance ⓘ

instance1 (e7697f66-be1b-4482-b6f9-2b4f8f9677e8) ▼

Device Name ⓘ

/dev/vdc

Cancel Attach Volume

6. Observe that the volume is attached to the desired instance.

<input type="checkbox"/>	Name	Description	Size	Status	Type	Attached To
<input type="checkbox"/>	volume1	Data volume	1GiB	In-use	standard	Attached to <a href="#">instance1</a> on /dev/vdb

7. Access *instance1* console (Project → Compute → Instances → Instance1 → Console) and login as root (password r3dh4t1!).
8. Run the volume-config.sh script (/usr/local/sbin/volume-config.sh)

```
[root@instance1 ~]# volume-config.sh
preparing data disk...
[
[

Type your name: OpenStack
done!
[root@instance1 ~]# _
```

This script was created for this Test Drive and does the following:

If unused disk:

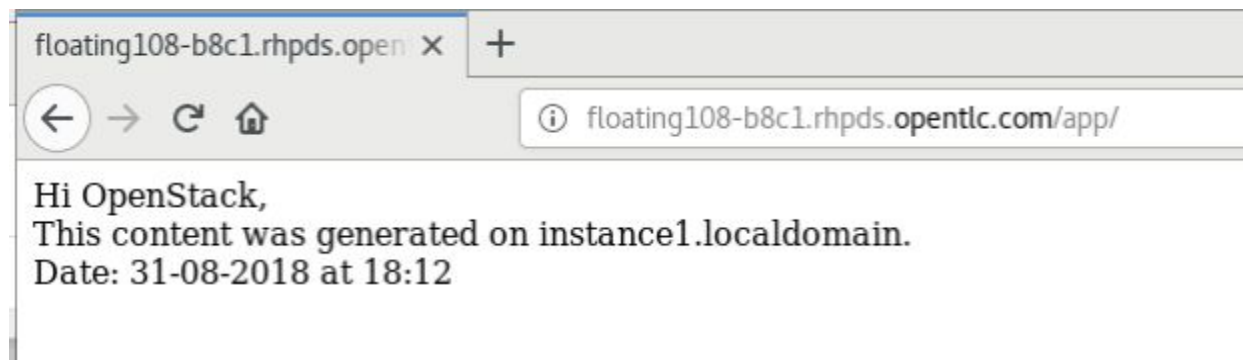
**i**

- Create partition
- Create filesystem
- Mount filesystem
- Create an HTML file based on user input

If it finds data in disk:

- Mount filesystem

9. Access *instance1* via HTTP. If instance1 has the floating IP 192.168.200.108, use the following URL: <http://floating108-GUID.rhpdn.opentlc.com/app>.



10. From OpenStack Dashboard, **delete instance1** (Project → Compute → Instances → instance1 Actions → Delete Instance).
11. From OpenStack Dashboard (Project → Compute → Volumes, Manage Attachments), attach *volume1* to *instance2*.

Manage Volume Attachments ✕

Instance	Device	Actions
No items to display.		

Attach To Instance

Attach to Instance <sup>?</sup> ⓘ

instance2 (11cfd4fe-bd97-4eb7-834c-6710113c20f8) ▼

Device Name ⓘ

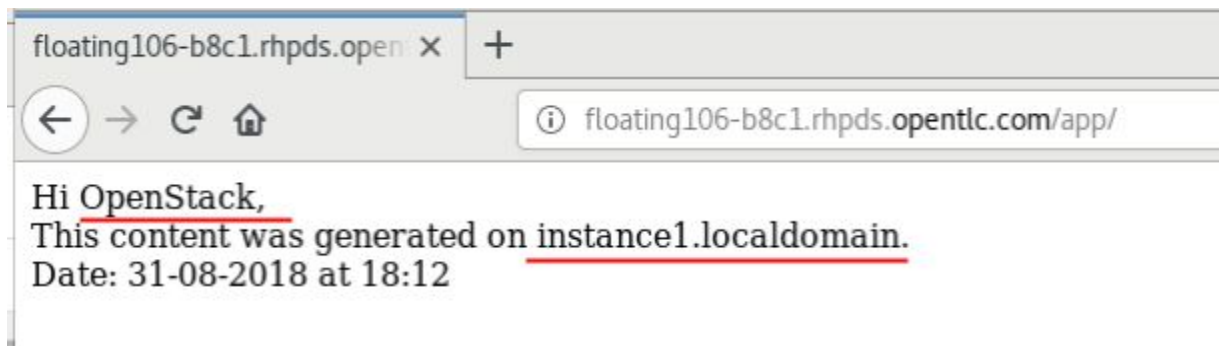
/dev/vdc

Cancel Attach Volume

12. Access *instance2* console (Project → Compute → Instances → Instance2 → Console) and login as root (password r3dh4t1!).
13. Run the volume-config.sh script (/usr/local/sbin/volume-config.sh)

```
[root@instance2 ~]# volume-config.sh
disk available. mounting partition...
[
done!
[root@instance2 ~]#
```

14. Access *instance2* via HTTP. If instance2 has the floating IP 192.168.200.106, use the following URL: <http://floating106-GUID.rhpbs.opentlc.com/app>.





## Lab 5: Key Pairs

In this lab you'll learn:

- What is a key pair.
- How to create key pairs in OpenStack.

### Overview of Key Pairs

Key pairs are SSH credentials based on public and private keys. Public keys are usually injected to instances when they are created, so that they can be accessed by whoever has the corresponding private key, eliminating the need to configure default passwords on images.

To enable key pair injection to the instances, the most common solution is to have the cloud-init package installed on images, which is true for Red Hat provided images.



Although it is possible to use cloud-init to set password for users, this isn't recommended as it's technically possible for any instance in the same network to retrieve this information.

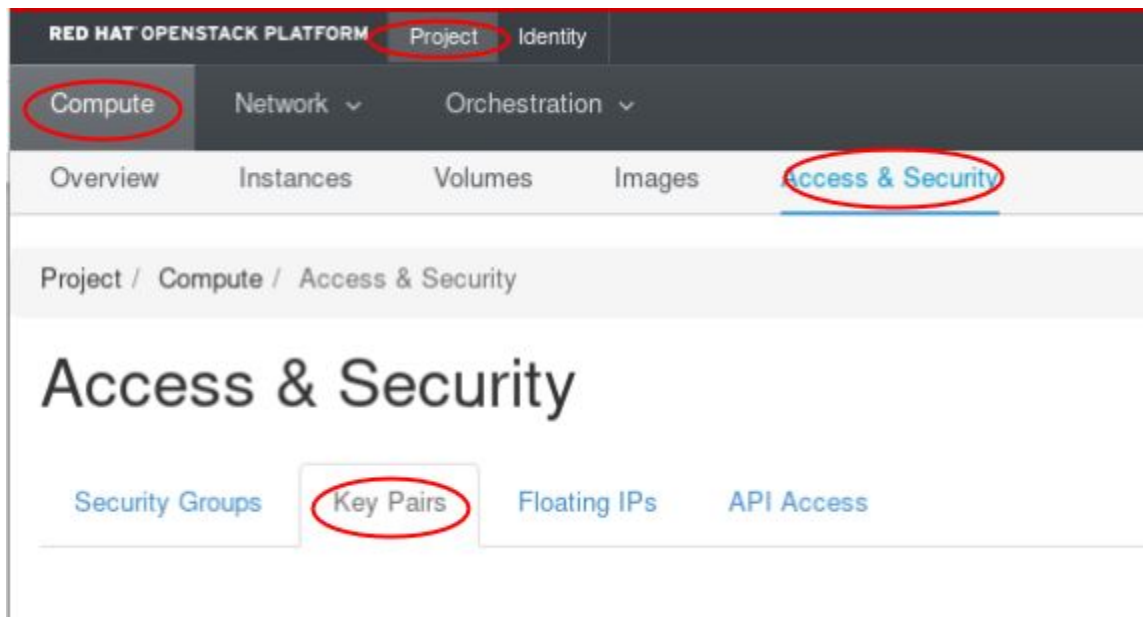
### Create a Key Pair

Although it's possible to upload an existing public key to OpenStack, in this lab you'll be creating a new key pair.

To create a key pair, use the following steps.

1. Log into the OpenStack Dashboard as *td-user*.

2. Go to Project → Compute → Access & Security → Key Pairs.



3. Click on “+Create Key Pair” button.



4. In the next dialog, enter “test-drive” as the name for the Key Pair and click on the Create Key Pair button.

A screenshot of the 'Create Key Pair' dialog box. The title is 'Create Key Pair' with a close button (X) in the top right corner. Below the title, there is a label 'Key Pair Name' followed by a text input field containing the text 'test-drive'. To the right of the input field, there is explanatory text: 'Key pairs are SSH credentials which are injected into images when they are launched. Creating a new key pair registers the public key and downloads the private key (a .pem file). Protect and use the key as you would any normal SSH private key.' At the bottom right of the dialog, there are two buttons: 'Cancel' and 'Create Key Pair'.

5. Proceed to save the *test-drive.pem* file in your local machine. This is the private key which should be kept secure.



The location where the private key file will be saved depends on your browser configuration.

## Use a Key Pair When Creating a RHEL instance

Creating an instance with an associated key pair is done the same way you did before with the exception that you need to ensure the key pair is selected in the corresponding dialog, as follows.

1. Log into the OpenStack Dashboard as *td-user*.
2. Go to Project → Compute → Instances, select *instance1* and *instance2* and **delete both** - this will automatically release the floating IPs.
3. Go to Project → Compute → Instances and click on the “Launch Instance” button.
4. Use the following information to create the instance.

Instance name	<b>instance3</b>
Availability Zone	nova
Count	1
Select Boot Source	Image
Create New Volume	No
Image	<b>RHEL-7</b>
Flavor	td.small
Network	private
Security Groups	default SSH_from_any
Key Pair	<b>test-drive</b>

5. After the instance is launched, associate floating IP **192.168.200.106** to the instance, selecting the corresponding option in the “Actions” menu.

## SSH into the instance

As *instance3* instance has a floating IP, any device that has access to network 192.168.200.0/24 should be able to reach this instance. Use the following steps to ssh into the instance.

1. From Linux ssh client:

```
$ ssh -i /path/to/test-drive.pem \  
cloud-user@floating106-GUID.rhps.opentlc.com
```

2. From Windows client with PuTTY:

- a. Download PuTTYgen from [here](#).
- b. Open PuTTYgen.
- c. Click Conversions, then click import key.
- d. Locate test-drive.pem key file, then click open.
- e. You can now save your key as a PPK file by clicking the Save private key button. Name the file **test-drive.ppk**.
- f. Use *floating106-GUID.rhps.opentlc.com* as the host to connect to.
- g. In Connection → SSH → Auth, click on *Browse* and select test-drive.ppk.
- h. Click on Open.
- i. Login name is *cloud-user*.

## Lab 6: Multi-tenancy

In this lab you'll learn:

- How the use of OpenStack Projects enable multi-tenancy

### Overview of Multi-tenancy

In a multi-tenant environment, resources like compute, networking and storage are isolated among different users. OpenStack implements this isolation through *projects*, allowing different organizations, customers or accounts, to use the same shared environment. This is managed by the Identity Service, known by the project name keystone.

Each project has its own configured quota for several kind of infrastructure resources, like vCPUs, RAM, disk space, number of persistent volumes, floating IP addresses, etc. This allows the administrator to restrict how much resources each project will use.

In order to have access to some resource, the user must have access to the project where that resource belongs, unless the resource is public or shared, like some networks or images.

Besides, users may have different roles. Default Red Hat OpenStack Platform roles follows:

admin	OpenStack administrator with access to all operations. Can manage projects and users.
_member_	Regular user that can create/delete instances, create/delete project networks, create/delete volumes on projects he has access to.
heat_stack_user	Regular user that can use heat orchestration service to manage resources.

Identity V3 included support to domains, which are high-level containers for projects, users and groups. This allows administrators to configure roles for the user at the project or at the domain level, enabling more fine grained Role Based Access Control (RBAC) capabilities. For example, it's possible to configure an administrator that can manage other users and projects for an specific domain.

It's common to use the terms *project* and *tenant* interchangeably.

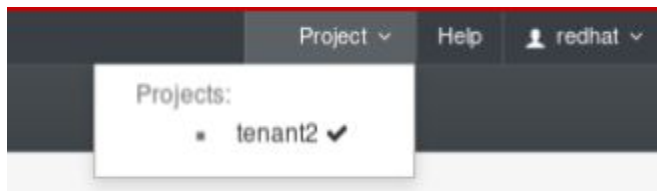
## Access a Different Project

In this environment there's a user which is associated to a project different than the "test-drive" being used so far.

1. Login using the following credentials:

Login	redhat
Password	r3dh4t1!

2. Notice on the top right corner that this user sees only a project called "tenant2".



3. Navigate to Project → Compute → Instances and start the existing *instance1*, which although may have the same name of instances of other projects, is a completely different one.
4. Access instance console and log into it using the username "cirros" and the password "cubswin:)", without the quotes. This instance is using a lightweight linux system called CirrOS, used mainly for tests like this.
5. Try pinging 8.8.8.8 or any other Internet public address and notice it won't work, as this project doesn't have a virtual router connecting the private network to the external one.
6. From OpenStack Dashboard, visualize Network Topology (Project → Network → Network Topology). Note how different this is from the "test-drive" tenant.

## Lab 7: Provisioning Instances using CloudForms

Red Hat® CloudForms® is an infrastructure management platform that allows IT departments to control users' self-service abilities to provision, manage, and ensure compliance across virtual machines and private clouds.

Provisioning refers to the capacity an infrastructure has to deliver a resource and manage its life cycle. Provisionable resources can include virtual machines, Instances, storage space, or any resource a given infrastructure can manage. The web interface on a *CloudForms* appliance provides an easy way for end users and administrators to provision new virtual machines and Instances. Virtual machines can be provisioned from virtual machine templates, PXE boot, or ISO images. Instances can be provisioned from images.

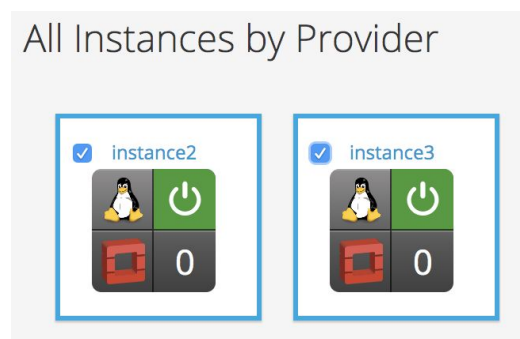
### Cleaning Up Our Environment

First, we need to clean up our environment so that we can reuse resources such as CPU, Memory, vDisks and of course ours Floating IPs.

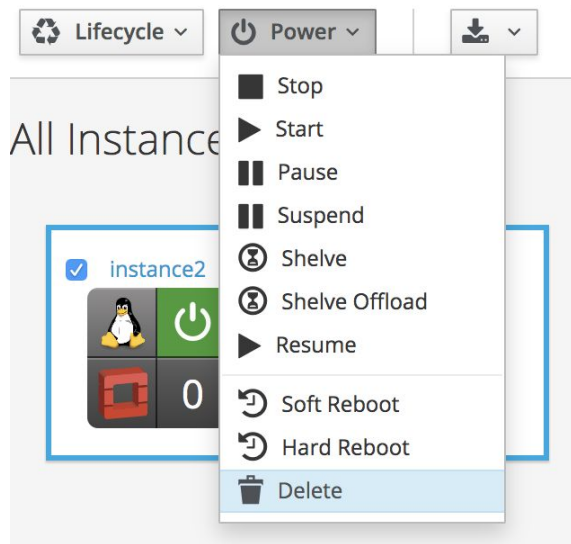
1. Access to **CloudForms Cloud Management Portal**, using the URL provided to you (eg. <https://cloudforms-GUID.rhpds.opentlc.com>)



Login	td-admin
Password	r3dh4t!

2. Go to: **Compute** → **Clouds** → **Instances** and select **instance2** and **instance3**:




Then click on **Power → Delete**:

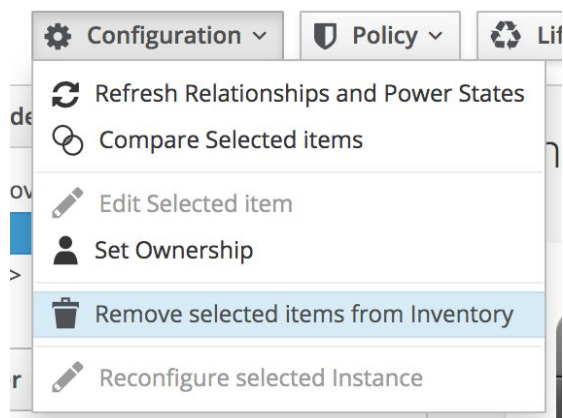


Click on **Refresh**  button until both instances are in **Terminate** state .

Select again **instance2** and **instance3**:




Then click on **Configuration → Remove selected items from inventory**:



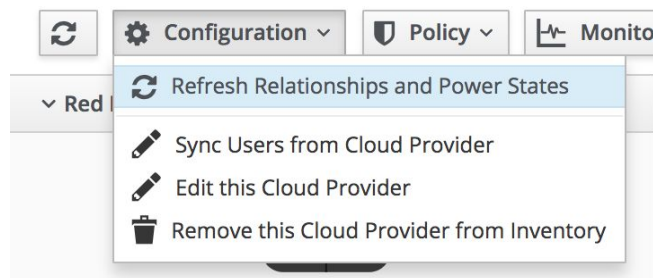
Click again on **Refresh**  button until both instances disappears.




Go to **Compute** → **Clouds** → **Providers** and then click on **Red Hat OpenStack Platform** provider:



After that, click on **Configuration** → **Refresh Relationships and Power States** to update all information about our Red Hat OpenStack Provider:



Finally, click on **Refresh**  button, until you can see “**Last Refresh: Success - Less Than A Minute Ago**” Status

### Red Hat OpenStack Platform (Summary)

Properties	
Region	OpenStack
Host Name	192.168.200.10
Type	OpenStack
API Port	13000
Management Engine GUID	640d6a41-823e-4124-bbee-19c657a3ab72
Region	regionOne

Status	
Default Credentials	Valid
Last Refresh	Success - Less Than A Minute Ago
Last Refresh Date	2018-08-31 20:28:22 -0500

**Now we're ready to next Labs!**

## Creating a New Instance Using CloudForms

1. Access to CloudForms Cloud Management Portal, using the URL provided to you (eg. <https://cloudforms-GUID.rhpds.opentlc.com>)

Login	td-admin
Password	r3dh4t!

2. Go to: **Compute** → **Clouds** → **Instances** and click on **Lifecycle** → **+Provision Instances**

Select the Template “**RHEL-7**” as Image. Click on **Continue** and fill out the information as follows:

### *Request Tab:*

Email	your.email@yourdomain.com
First/Last Name	your name (if you want. Not necessary)

### *Purpose Tab:*

Purpose	Select <b>Environment</b> → <b>Test</b>
---------	---

### *Catalog Tab:*

Number of instances	1
Instance Name	wordpress-instance

### Environment Tab:

- Set **Placement - Options** as follows:

Cloud Tenant	test-drive
Availability Zones	nova
Cloud Network	private
Security Groups	SSH_HTTP_From_Any
Public IP Address	192.168.200.106

Request

Purpose

Catalog

Environment

Prop

Placement

Choose Automatically

☐

Placement - Options

Cloud Tenant

test-drive

▼

Availability Zones

nova

▼

Cloud Network \*

private

▼

Security Groups

SSH\_HTTP\_From\_Any

▼

Public IP Address

192.168.200.106

▼

### Properties Tab:

Instance Type	m1.small
Guest Access Key Pair	test-drive

Request
Purpose
Catalog
Environment
Properties

Properties

Instance Type \*

m1.small: 1 CPUs, 2 GB RAM, 20 (v

Guest Access Key Pair

test-drive v

**Schedule Tab:**

Time until Retirement	1 Month
Retirement Warning	1 Week

3. Click on **Submit** to create the new Instance.

4. Switch to **OpenStack Dashboard**, using the URL provided to you (eg. **https://openstack-GUID.rhpds.opentlc.com**)

Observe the deployment of the new OpenStack instance. Click on **Instances** → **wordpress-instance** to see the instance details:

Login	td-admin
Password	r3dh4t!

5. Return to **CloudForms Management Portal**. Go to **Services** → **Requests**. Click on **Reload** button, until the **Last Message** “**CF\_Worker1: VM [wordpress-instance] Provisioned Successfully**” appears and **Request State** is in **Finished** state:

Last Message				Status	Request State	Request ID	Requester
[CF_Worker1] VM [wordpress-instance] IP [10.10.10.112] Provisioned Successfully		<input type="checkbox"/>	<input checked="" type="checkbox"/>	Ok	Finished	18	Test-Drive Admin

## Lab 8: Configuring Instances Automatically using Ansible Tower

Ansible seamlessly unites workflow orchestration with configuration management, provisioning, and application deployment in one easy-to-use and deploy platform.

Regardless of where you start with Ansible, you'll find our simple, powerful and agentless automation platform has the capabilities to solve your most challenging problems.

Centralizing configuration file management and deployment is a common use case for Ansible, and it's how many power users are first introduced to the Ansible automation platform. In that way, it's very easy to automate the configuration of your bare-metal servers, virtual machines and cloud instances.

When you define your application with Ansible, and manage the deployment with Ansible Tower, teams are able to effectively manage the entire application lifecycle from development to production.


Access to **Ansible Tower Portal**, using the URL provided to you (eg. <https://tower-GUID.rhpdsopentlc.com>)

Login	td-admin
Password	r3dh4t1!

### Setting Up Credentials in Ansible Tower

Credentials are utilized by Tower for authentication when launching Jobs against machines, synchronizing with inventory sources, and importing project content from a version control system.

You can grant users and teams the ability to use these credentials, without actually exposing the credential to the user. If you have a user move to a different team or leave the organization, you don't have to re-key all of your systems just because that credential was available in Tower.

Navigate within the **Ansible Tower Portal** and click on **Credentials** (within **RESOURCES** Section). Click on  to add new Credentials:

Create the new Credentials as follows:

NAME	Machine Credentials
DESCRIPTION	Machine Credentials for OSP Instances
ORGANIZATION	Default
TYPE	Machine
USERNAME	cloud-user
SSH PRIVATE KEY	Drag and drop your <b>test-drive.pem</b> file
PRIVILEGE ESCALATION METHOD	sudo

Click on **SAVE** to save the changes.

Machine Credentials

DETAILS

PERMISSIONS

\* NAME ?

Machine Credentials

DESCRIPTION ?

Machine Credentials for OSP Instances

ORGANIZATION

Q Default

\* CREDENTIAL TYPE ?

Q Machine

TYPE DETAILS

USUARIO

cloud-user

CONTRASEÑA

SHOW

☐ Prompt on launch

SSH PRIVATE KEY HINT: Drag and drop an SSH private key file on the field below.

REVERT

JKuPEpwkI3RJbCyMv17MC9UBAoGAJneg+mkejDaR9NYoqUzdevhL9FYkLdDXV8zh  
ujmZfLL1SofN86m9isgpEuXBHM9100eS2YK83V6zMz+fstI4IjjqQxjYqACcm6z  
AGVI4vXV7e/vjngxNAMI5VdF0eyNSHpufmF0iR1aD+xDAC/QVg6hoEq28mXHUBs  
1uIP1TUCgYABlr+HpEKKeDeOR30eC/fkDTHX19B5nQfM13X3N4eJ5n8KDZectdxk  
I82XdPbf7M54N6YJN56vbruj3aNCdj+IucAVYlqvL51LvIaQ087ASKGBvIW0I28k  
lQ/jTnrYzU58eP1/ZPD8xNSKNECRs550th70nounY7Y+dzo8wLXGEg==  
-----END RSA PRIVATE KEY-----

PRIVATE KEY PASSPHRASE

☐ Prompt on launch

PRIVILEGE ESCALATION METHOD ?

PRIVILEGE ESCAL

SHOW


sudo

# Creating Ansible Tower Job Templates

A job template is a definition and set of parameters for running an Ansible job. Job templates are useful to execute the same job many times. Job templates also encourage the reuse of Ansible playbook content and collaboration between teams. While the REST API allows for the execution of jobs directly, Tower requires that you first create a job template.

Navigate within the **Ansible Tower Portal** and click on **TEMPLATES**.


This menu opens a list of the job templates that are currently available. The Templates tab also enables the user to launch, schedule, modify, and remove a job template.

1. Click on **Templates** (within **RESOURCES** Section) and then click on  button (selecting **Job Template**) to create a new Job Template to deploy a **Wordpress Application** within our **wordpress-instance**:

NAME	Deploy Wordpress App on OSP Instances
DESCRIPTION	Deploy Wordpress App on OSP Instances
JOB TYPE	Run
INVENTORY	OpenStack Inventory
PROJECT	OpenStack Project
PLAYBOOK	osp-tower-ansible-install-wordpress.yml
CREDENTIAL	Machine Credentials
LIMIT	wordpress-instance
VERBOSITY	0 (Normal)
OPTIONS	Enable Privilege Escalation: <b>checked</b>


Click on **SAVE** to save the changes.

**TIP: Selecting a Credential:** just click on  icon and then click on **Machine Credentials**. Finally click on **SELECT** button:

**CREDENTIALS** 

SELECTED: ✕ MACHINE: Machine Credentials REVERT

CREDENTIAL TYPE: Machine ▼

SEARCH  KEY

NAME ▲

☐ Demo Credential

☐ Localhost credentials

☒ Machine Credentials

ITEMS 1 - 3

CANCEL SELECT


Your new Job Template now looks as follows:

**Deploy Wordpress App on OSP Instances**

DETAILS PERMISSIONS NOTIFICATIONS COMPLETED JOBS ADD SURVEY


\* NAME

Deploy Wordpress App on OSP Instances

\* INVENTORY 

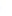
☐ PROMPT ON LAUNCH

Q OpenStack Inventory

\* CREDENTIAL 


☐ PROMPT ON LAUNCH

Q ✕ MACHINE: Machine Credentials

\* VERBOSITY 


☐ PROMPT ON LAUNCH


0 (Normal)

SKIP TAGS 

☐ PROMPT ON LAUNCH


OPTIONS

☒ Enable Privilege Escalation 


☐ Allow Provisioning Callbacks 

DESCRIPTION


Deploy Wordpress App on OSP Instances

\* PROJECT 


Q OpenStack Project

FORKS 

DEFAULT

INSTANCE GROUPS 

Q

LABELS 

\* JOB TYPE 

Ejecutar

\* PLAYBOOK 

osp-tower-ansible-install-wordpress.yml

LIMIT 


wordpress-instance

JOB TAGS 

SHOW CHANGES 

OFF



2. Click again on **Templates** (within **RESOURCES** Section) and then click on  button (selecting **Job Template**) to create a new Job Template to deploy a **Ticketmonster Web Application** within an instance. We'll be creating this instance via *CloudForms* using a Service Bundle and *we'll be using this Job Template later within this Service Bundle to deploy the application within the instance:*

NAME	Deploy Ticketmonster App on OSP Instances
DESCRIPTION	Deploy Ticketmonster App on OSP Instances
JOB TYPE	Run
INVENTORY	OpenStack Inventory
PROJECT	OpenStack Project
PLAYBOOK	osp-tower-ansible-install-ticketmonster.yml
CREDENTIAL	Machine Credentials
LIMIT	<b>Keep it in blank</b> (just → Prompt on Launch: <b>checked</b> )
VERBOSITY	0 (Normal)
OPTIONS	Enable Privilege Escalation: <b>checked</b>

Click on **SAVE** to save the changes.

Deploy Ticketmonster App on OSP Instances

DETAILS PERMISSIONS NOTIFICATIONS COMPLETED JOBS ADD SURVEY

\* NAME Deploy Ticketmonster App on OSP Instances

DESCRIPTION Deploy Ticketmonster App on OSP Instances

\* JOB TYPE Ejecutar ☐ PROMPT ON LAUNCH

\* INVENTORY OpenStack Inventory ☐ PROMPT ON LAUNCH

\* PROJECT OpenStack Project

\* PLAYBOOK osp-tower-ansible-install-ticketmonster.yml ☒ PROMPT ON LAUNCH

\* CREDENTIAL Machine Credentials ☐ PROMPT ON LAUNCH

FORKS DEFAULT

LIMIT ☐ PROMPT ON LAUNCH

\* VERBOSITY 0 (Normal) ☐ PROMPT ON LAUNCH

INSTANCE GROUPS

JOB TAGS ☐ PROMPT ON LAUNCH

SKIP TAGS ☐ PROMPT ON LAUNCH

LABELS

SHOW CHANGES ☐ PROMPT ON LAUNCH

OFF

OPTIONS

☒ Enable Privilege Escalation

## Launching Ansible Tower Job Templates

Now it's time to deploy our new **Wordpress** App within **wordpress-instance** previously created via CloudForms.

First, use the URL provided to you to be sure **Wordpress** is not installed, using it with a web browser:

**<http://floating106-GUID.rhpds.opentlc.com/wordpress>**

Then, go back to **Ansible Tower Portal**. Click on **Templates** (within **RESOURCES** Section) and find the “**Deploy Wordpress App on OSP Instances**” Job Template, using the **SEARCH** textbox as follows:

TEMPLATES **1**

wordpress



KEY

x wordpress CLEAR ALL


NAME ^

TYPE ^

ACTIVITY

Deploy Wordpress App on OSP Instances

Job Template

Click on  to start a job creation using this template.

Start a job using this template

ACTIONS

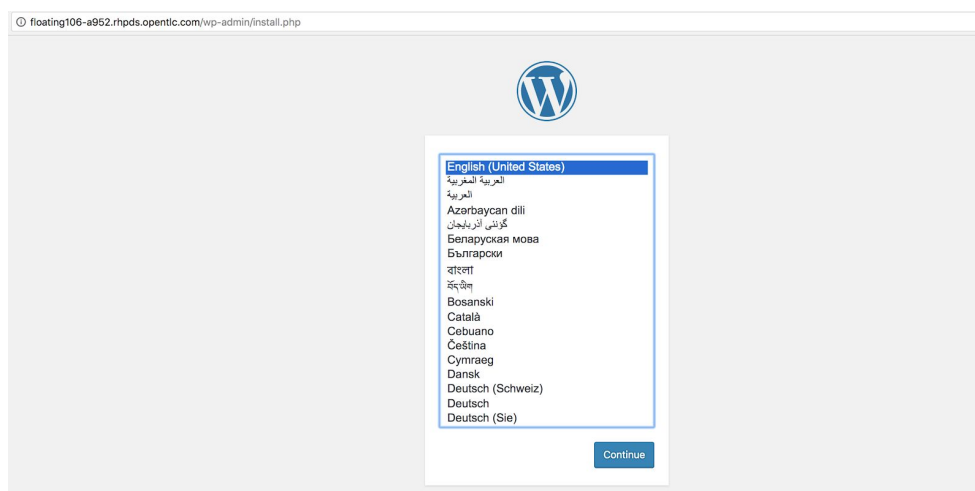


This Action will launch the “Deploy Wordpress App on OSP Instances” Job Template, starting the Website installation. You will be redirected to a dynamic Job Output page:

The screenshot displays the Ansible Tower interface. On the left, the 'DETAILS' panel shows job information: STATUS is 'Successful', STARTED is '8/31/2018 10:24:21 PM', FINISHED is '8/31/2018 10:26:16 PM', TEMPLATE is 'Deploy Wordpress App on OSP Instances', JOB TYPE is 'Run', LAUNCHED BY is 'td-admin', INVENTORY is 'OpenStack Inventory', PROJECT is 'OpenStack Project', REVISION is '902bcd5', PLAYBOOK is 'osp-tower-ansible-install-wordpress.yml', MACHINE CREDENTIAL is 'Machine Credentials', FORKS is '0', LIMIT is 'wordpress-instance', VERBOSITY is '0 (Normal)', and INSTANCE GROUP is 'tower'. On the right, the 'Deploy Wordpress App on OSP Instances' job output is shown. It includes a search bar and a list of log entries. Key entries include: 'changed: [d181e588-1a43-4c2e-93c8-9e13decfe54f]', 'RUNNING HANDLER [mariadb : restart mariadb]', 'RUNNING HANDLER [nginx : restart nginx]', 'RUNNING HANDLER [php-fpm : restart php-fpm]', and 'PLAY RECAP d181e588-1a43-4c2e-93c8-9e13decfe54f : ok=46 changed=44 unreachable=0 failed=0'.

You will see the Status of the Job, Started / Finished Time, Tasks, etc.

If your Job finished **Ok**, with Status **Successful**, go to your **Wordpress URL** again and make sure your Wordpress Server is **Up & Running**. Play a moment with your new Wordpress!



You can take a look at the **Ansible Playbook** we are using for this Job Template:

<https://github.com/rcalvaga/openstack-ansible/blob/master/osp-tower-ansible-install-wordpress.yml>

# Lab 9: Automating Services Delivery using CloudForms and Ansible Tower

Your apps have to live somewhere. If you're PXE booting and kickstarting bare-metal servers or VMs, or creating virtual or cloud instances from templates, Ansible and Red Hat® Ansible® Tower help streamline the process in conjunction with Red Hat® CloudForms®.

We are using *Ansible Tower* as an Automation Provider for *CloudForms*, and we want to successfully provision a Web Application (Ticketmonster Service Portal) within an OpenStack Instance. For this task, we need to create a new **Service Bundle** based on an OpenStack Instance Service Item to create a new Instance + Ansible Tower Job Template Service Item to deploy our Ticketmonster Web Application within it, and then access to *CloudForms Self-Service Portal* and select our new service to build our own Application.

## 1. Access to CloudForms Cloud Management Portal,

Login	td-admin
Password	r3dh4t1!

## Creating Service Catalog Items and Bundles

Through the use of catalogs, Red Hat CloudForms provides support for multi-tier service provisioning to deploy layered workloads across hybrid environments. You can create customized dialogs that will give consumers of the services the ability to input just a few parameters and provision the entire service. The following table lists the terminology associated with catalogs that you will use within the CloudForms user interface for service provisioning.

Now we need to create first two Service Catalog Items:

- **Create OpenStack Instance.** Based on Catalog Item Type: OpenStack.
- **Deploy Ticketmonster App on OSP Instance.** Based on Catalog Item Type: AnsibleTower.

Later, we'll be creating a Service Catalog Bundle using this Service Catalog Items.

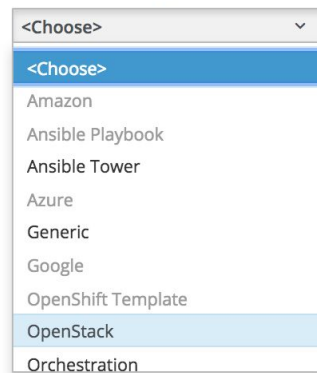
## A. Create an OpenStack Service Catalog Item:

1. Navigate to the **Catalogs** section in the accordion, **Services → Catalogs** and then **Catalog Items → All Catalog Items → OpenStack Catalog** and click on **Configuration → Add a New Catalog Item**.

Select **OpenStack** from the drop-down list:

Adding a new Service Catalog Item

Catalog Item Type



A screenshot of a web application's dropdown menu. The dropdown is titled '<Choose>' and lists several options: '<Choose>', 'Amazon', 'Ansible Playbook', 'Ansible Tower', 'Azure', 'Generic', 'Google', 'OpenShift Template', 'OpenStack', and 'Orchestration'. The 'OpenStack' option is currently selected and highlighted in blue.

Fill out information as follows:

### Basic Info:

Name/Description	Create OSP Instance
Display in Catalog	<b>Keep it unchecked</b>
Catalog	OpenStack Catalog
Dialog	<b>&lt;No Dialog&gt;</b>

### Request Info:

Click on **Request Info** Tab and then:

### Catalog Tab:

Select Image Name	RHEL-7
-------------------	--------

Number of Instances	1
Instance Name	changeme

#### Environment Tab:

Cloud Tenant	test-drive
Availability Zones	nova
Cloud Network	private
Security Groups	SSH_HTTP_From_Any
Public IP Address	<None> (we will select it using a Service Dialog later)

#### Properties Tab:

Instance Type	m1.small
Guest Access Key Pair	test-drive

Request
Purpose
Catalog
Environment
Properties

Properties

Instance Type \*

m1.small: 1 CPUs, 2 GB RAM, 20 (v

Guest Access Key Pair

test-drive v

2. Click on **Add** to create the new OSP Service Item.

## B. Create an Ansible Tower Service Catalog Item:

1. Navigate to the **Catalogs** section in the accordion, **Services → Catalogs** and then **Catalog Items → All Catalog Items → OpenStack Catalog** and click on **Configuration → Add a New Catalog Item**.

Select **AnsibleTower** from the drop-down list:

Adding a new Service Catalog Item

Catalog Item Type

<Choose>

<Choose>

Amazon

Ansible Playbook

Ansible Tower

Fill out information as follows:









### Basic Info:

Name/Description	Deploy Ticketmonster App on OSP Instance
Display in Catalog	<b>Keep it unchecked</b>
Catalog	OpenStack Catalog
Dialog	<b>&lt;No Dialog&gt;</b>
Provider	Ansible Tower 3 Automation Manager
Ansible Tower Job Template	Deploy Ticketmonster App on OSP Instances
Provisioning Entry Point	<b>/AutomationManagement/AnsibleTower/Service/Provisioning/StateMachines/Provision/provision_from_bundle</b>

**TIP:** Selecting the Provisioning Entry Point Path:

Click on the **Provisioning Entry Point Path** and then select it as follows:

### Select Entry Point Instance

- ▼  Datastore
  - ▼  Test-Drive
    - ▼  AutomationManagement
      - ▼  AnsibleTower
        - ▼  Service
          - ▼  Provisioning
            - ▼  StateMachines
              - ▼  Provision

 provision\_from\_bundle

Click on **Apply** to select the Provisioning Entry Point

Then click on **Add** to create our new AnsibleTower Service Item.

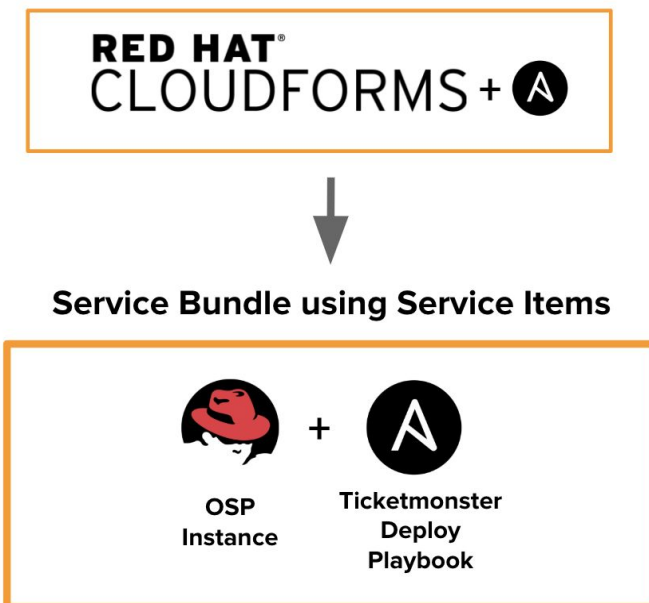


## Adding a new Service Catalog Item

Name / Description	Deploy Ticketmonster App on OSF	Deploy Ticketmonster App on OSF	<input type="checkbox"/> Display in Catalog
Catalog	OpenStack Catalog		
Dialog	<No Dialog>		
Provider	Ansible Tower 3 Automation Mai		
Ansible Tower Job Template	Deploy Ticketmonster App on OSF		
Provisioning Entry Point	/AutomationManagement/AnsibleTower/Service/Provisioning/StateMachines/Provision/provision_from_bundle		
Reconfigure Entry Point			
Retirement Entry Point			

### C. Create a Ticketmonster App Service Catalog Bundle:

Now it's time to create a Service Catalog Bundle, using our Service Catalog Items created previously, so that we can create first an OpenStack Instance and then automate a configuration and an application deployment within it.



1. Navigate to the **Catalogs** section in the accordion, **Services** → **Catalogs** and then **Catalog Items** → **All Catalog Items** → **OpenStack Catalog** and click on **Configuration** → **Add a New Catalog Bundle**.

Fill out information as follows:

**Basic Info:**

Name/Description	Ticketmonster Service
Display in Catalog	<b>Checked</b>
Catalog	OpenStack Catalog
Dialog	<b>Ticketmonster Service Dialog</b>

**Details:**

Click on **Details Tab** and put “**My JBoss Service**” as Long Description:

Basic Info

Details

Resources

Long Description

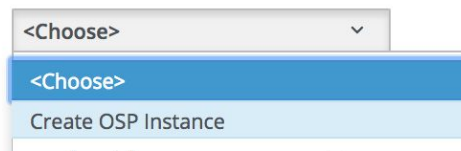
1 My JBoss Service

**Resources:**

Click on **Add a Resource** Dropdown and Select: “**Create OSP Instance**”:

## Resources

Add a Resource



<Choose> ▾

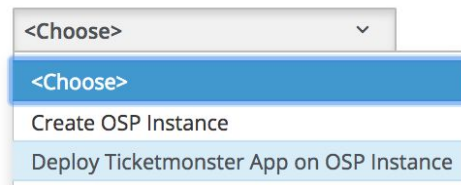
<Choose>

Create OSP Instance

Click again on **Add a Resource** Dropdown and Select: **“Deploy Ticketmonster App on OSP Instance”**:

## Resources

Add a Resource



<Choose> ▾

<Choose>

Create OSP Instance

Deploy Ticketmonster App on OSP Instance

After that, it's necessary to adjust some stuff for our selected resources, such as **Action Order**, **Provision Order** and **Delay** as follows:

Name	Create OSP Instance
Action Order	1
Provision Order	1

Name	Deploy Ticketmonster App on OSP Instance
Action Order	2
Provision Order	2
Delay (mins) Start	1

You will have something like that:

Selected Resources

	Name	Action Order	Provision Order	Action		Delay (mins)	
				Start	Stop	Start	Stop
+	Create OSP Instance	1 ↕	1 ↕	Power On ↕	Shutdown ↕	None ↕	None ↕
+	Deploy Ticketmonster App on OSP Instance	2 ↕	2 ↕	Power On ↕	Shutdown ↕	1 ↕	None ↕

Then click on **Add** to create the new Service Bundle.

Now it's time to order our new Service Bundle!

## Ordering a new Service using CloudForms Self-Service Portal

Red Hat CloudForms Self Service is a web-based graphical user interface for ordering and managing IT service requests. You can enable self-service tenant end users, who can easily access their services, track requests, and manage their accounts using the Self Service user interface (SSUI), which has widgets, dashboard controls and feedback. The Self Service user interface supports role-based access control (RBAC) of menus and features, similar to in the full administrative user interface.

1. Access to **CloudForms Self-Service Portal**, using the URL provided to you, adding **/self\_service** context-root

(eg. [https://cloudforms-GUID.rhps.opentlc.com/self\\_service](https://cloudforms-GUID.rhps.opentlc.com/self_service))

Login	td-admin
Password	r3dh4t!

2. Click on **Service Catalog**

3. Click on **Ticketmonster Service**

Fill out required information as follows:

Service Name	My JBoss Service
VM Name	jboss-instance
Flavor	m1.small
Floating IP	192.168.200.108

**Ticketmonster Service**

**New section**

Service Name

My JBoss Service

VM Name

jboss-instance

Flavor

m1.small on Red Hat OpenStack Platform ▾

Floating IP

192.168.200.108 on Red Hat OpenStack Platform Networ ▾

4. Click on **Add to Shopping Cart** button. Then click on **your Shopping Cart** icon and click on **Order** to order your service.

Meanwhile, use the URL provided to you to be sure **Ticketmonster Web App** is not installed, using it with a web browser and adding **:8080** port and context-root **/ticket-monster** as follows:

**<http://floating108-GUID.rhpds.opentlc.com:8080/ticket-monster>**

**5. Return to the CloudForms Cloud Management Engine Portal:**

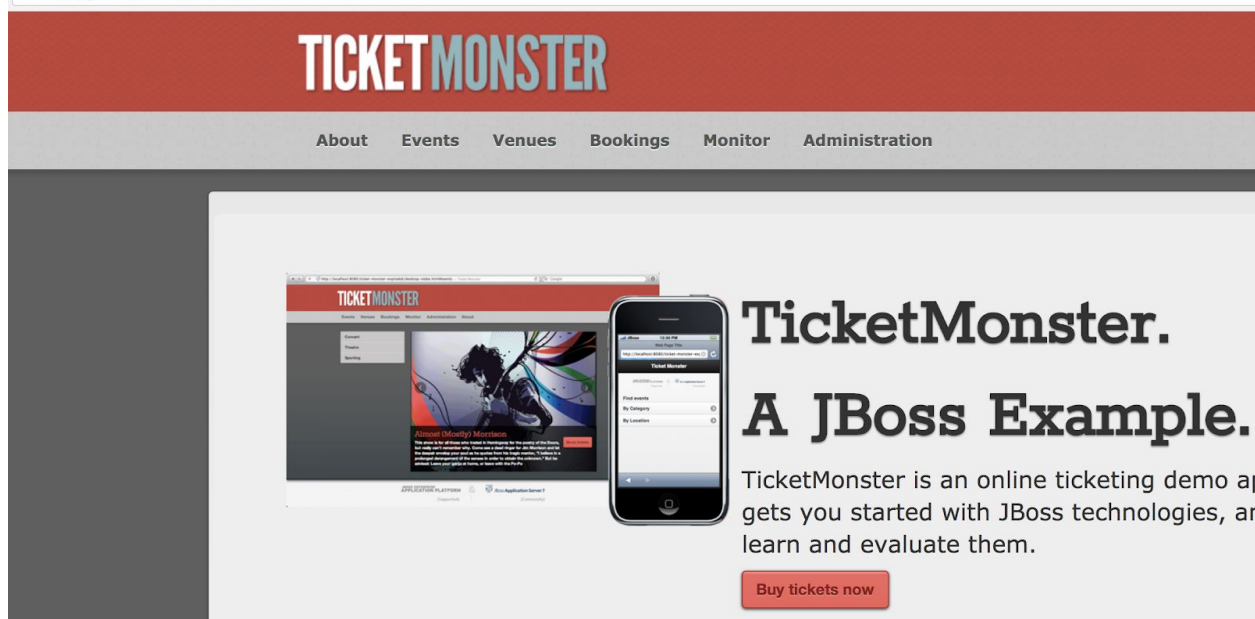
Go to **Services → Requests** and click on **Reload** button until Request State is in **Finished** state.

**6. Switch to RHEL OpenStack Platform Portal** to observe the deployment of the new OpenStack instance from a Service. Click on **Instances → jboss-instance** to see the instance details.

**7. You can access to Ansible Tower Portal.** Click on **Jobs**. Click on **Deploy Ticketmonster App on OSP Instances** Job and take a look at all information.

**8. Go to your URL provided but adding **:8080** port and context-root **/ticket-monster**:**  
**<http://floating108-GUID.rhpds.opentlc.com:8080/ticket-monster>**

floating108-a952.rhpd.s.opentlc.com:8080/ticket-monster/



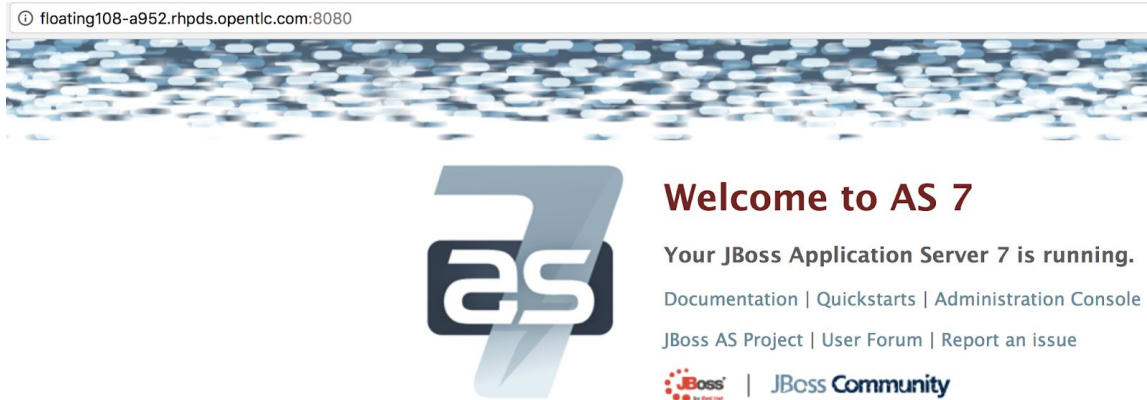
Now you can access your new **Ticketmonster Web App** using your link and it's time to play with it!

Go to your **JBoss Application Server URL** and make sure your App Server is Up & Running, accessing via **:8080 Port!**

**http://floating108-GUID.rhpd.s.opentlc.com:8080**

Click on **Administration Console** link. You can access the **JBoss Administration Console**:

**User:** admin / **Password:** Redhat1!



Click on **Deployments** → **Manage Deployments**. Refresh page if necessary. You will see your new deployments:

The image shows the JBoss AS 7.1 Administration Console interface. The left sidebar contains a navigation menu with sections like 'Server Status', 'Subsystem Metrics', 'Runtime Operations', and 'Deployments'. The 'Deployments' section is expanded, showing 'Manage Deployments'. The main content area is titled 'Deployments' and contains a table with the following data:

Name	Runtime Name	Enabled	En/Disable	Remove
helloworld.war	helloworld.war	✓	<button>Disable</button>	<button>Remove</button>
ticket-monster.war	ticket-monster.war	✓	<button>Disable</button>	<button>Remove</button>

At the bottom of the table, there are pagination controls showing '1-2 of 2'.

## TEST-DRIVE SURVEY:

If you finished the Test-Drive, please answer this Survey, selecting “**Private Cloud**” as **Test Drive Name**:

[https://docs.google.com/forms/d/e/1FAIpQLSdauHtguNMYICRE5x1nrE0Y11ASfDNnptSEqbLZi\\_TCsNgb2q/viewform](https://docs.google.com/forms/d/e/1FAIpQLSdauHtguNMYICRE5x1nrE0Y11ASfDNnptSEqbLZi_TCsNgb2q/viewform)